

Lab 2 - Data wrangling

Ayden Frost

2026-02-12

```
library(tidyverse)
```

Questions

Part 1

Question 1

```
midwest |>
  count(state, sort = TRUE)
```

```
# A tibble: 5 x 2
  state     n
  <chr> <int>
1 IL        102
2 IN         92
3 OH         88
4 MI         83
5 WI         72
```

Illinois has the highest number of counties with 102, while Wisconsin has the lowest number of counties with only 72.

Question 2

```
midwest |>
  count(county, state) |>
  count(county, name = "nstates") |>
  filter(nstates == 5)
```

```
# A tibble: 3 x 2
  county    nstates
  <chr>      <int>
1 CRAWFORD     5
2 JACKSON      5
3 MONROE       5
```

Question 3

```
midwest |>
  filter(popdensity > 25000) |>
  select(county, state, popdensity, poptotal, area) |>
  arrange(desc(popdensity))
```

```
# A tibble: 9 x 5
  county    state popdensity poptotal   area
  <chr>     <chr>     <dbl>      <int>   <dbl>
1 COOK      IL        88018.  5105067 0.058
2 MILWAUKEE WI        63952.  959275  0.015
3 WAYNE     MI        60334.  2111687 0.035
4 CUYAHOGA OH        54313.  1412140 0.026
5 DU PAGE   IL        39083.  781666  0.02
6 MARION    IN        34659.  797159  0.023
7 HAMILTON OH        34649.  866228  0.025
8 FRANKLIN OH        28278.  961437  0.034
9 MACOMB    MI        25621.  717400  0.028
```

```
midwest |>
  filter(popdensity == max(popdensity)) |>
  select(county, state, popdensity, poptotal, area)
```

```
# A tibble: 1 x 5
  county state popdensity poptotal   area
  <chr>  <chr>     <dbl>      <int>   <dbl>
1 COOK    IL        88018.  5105067 0.058
```

Question 4

```
midwest |>
  summarize(
    median(popdensity),
    q1 = quantile(popdensity, 0.25),
    q3 = quantile(popdensity, 0.75)
  )

# A tibble: 1 x 3
`median(popdensity)`    q1     q3
<dbl> <dbl> <dbl>
1       1156.   622.  2330
```

The distribution of population density of counties is unimodal and extremely right-skewed. A typical Midwestern county has population density of 1156 people per unit area. The middle 50% of the counties have population densities between 622 to 2330 people per unit area.

Question 5

```
midwest |>
  count(state, inmetro) |>
  group_by(state) |>
  mutate(prop = n / sum(n))
```

```
# A tibble: 10 x 4
# Groups:   state [5]
  state inmetro     n   prop
  <chr>    <int> <int> <dbl>
1 IL          0     74  0.725
2 IL          1     28  0.275
3 IN          0     55  0.598
4 IN          1     37  0.402
5 MI          0     58  0.699
6 MI          1     25  0.301
7 OH          0     48  0.545
8 OH          1     40  0.455
9 WI          0     52  0.722
10 WI         1     20  0.278
```

Question 6

```
midwest |>
  filter(percbelowpoverty >= 40,
         percollege <= 10) |>
  select(county, state,
         percbelowpoverty,
         percollege)

# A tibble: 1 x 4
  county    state percbelowpoverty percollege
  <chr>     <chr>        <dbl>       <dbl>
1 MENOMINEE WI            48.7        7.34
```

```
midwest |>
  filter(percollege >= 40,
         percbelowpoverty <= 20) |>
  select(county, state,
         percbelowpoverty,
         percollege)
```

```
# A tibble: 5 x 4
  county    state percbelowpoverty percollege
  <chr>     <chr>        <dbl>       <dbl>
1 CHAMPAIGN IL          15.6        41.3
2 DU PAGE   IL          2.71        42.8
3 HAMILTON IN          3.59        42.1
4 WASHTEAW MI          12.2        48.1
5 DANE      WI          10.5        43.6
```

```
midwest |>
  filter(
    (percbelowpoverty >= 40 & percollege <= 10) |
    (percbelowpoverty <= 20 & percollege >= 40)
  ) |>
  select(county, state,
         percbelowpoverty,
         percollege)
```

```
# A tibble: 6 x 4
  county     state percbelowpoverty percollege
  <chr>      <chr>        <dbl>       <dbl>
1 CHAMPAIGN IL          15.6        41.3
2 DU PAGE   IL          2.71        42.8
3 HAMILTON IN          3.59        42.1
4 WASHTENAW MI          12.2        48.1
5 DANE       WI          10.5        43.6
6 MENOMINEE WI          48.7        7.34
```

```

midwest |>
  mutate(
    potential_outlier = if_else(
      (percbelowpoverty >= 40 & percollege <= 10) |
      (percbelowpoverty <= 20 & percollege >= 40),
      "Yes",
      "No"
    )
  ) |>
  select(county, state,
         percbelowpoverty,
         percollege,
         potential_outlier) |>
  arrange(potential_outlier)

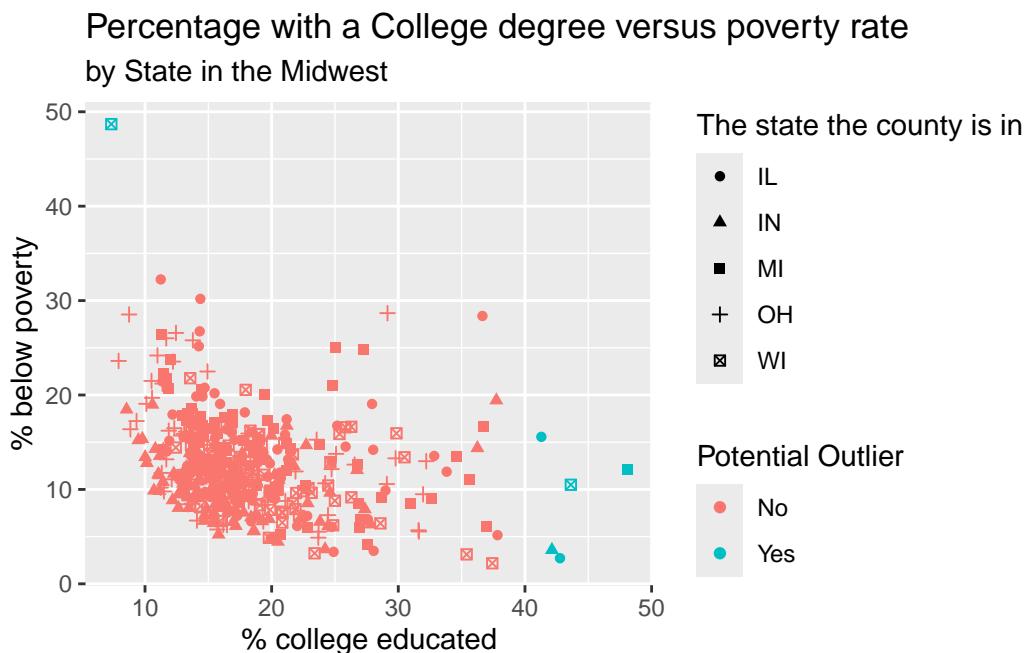
# A tibble: 437 x 5
  county   state percbelowpoverty percollege potential_outlier
  <chr>    <chr>        <dbl>       <dbl> <chr>
1 ADAMS    IL          13.2        19.6  No
2 ALEXANDER IL          32.2        11.2  No
3 BOND     IL          12.1        17.0  No
4 BOONE    IL          7.21        17.3  No
5 BROWN    IL          13.5        14.5  No
6 BUREAU   IL          10.4        18.9  No
7 CALHOUN  IL          15.1        11.9  No
8 CARROLL  IL          11.7        16.2  No
9 CASS     IL          13.9        14.1  No
10 CHRISTIAN IL         11.7        13.6  No
# i 427 more rows

```

```

midwest |>
  mutate(
    potential_outlier = if_else (
      (percbelowpoverty >= 40 & percollege <= 10) |
      (percbelowpoverty <= 20 & percollege >= 40),
      "Yes",
      "No"))
  ) |>
ggplot (aes (x = percollege, y = percbelowpoverty, colour = potential_outlier, shape = state),
geom_point() +
  labs (
    x= "% college educated",
    y = "% below poverty",
    colour = "Potential Outlier",
    shape = "The state the county is in",
    title = "Percentage with a College degree versus poverty rate",
    subtitle = "by State in the Midwest")

```



Question 7

```
state_population <- midwest |>
  group_by(state) |>
  summarize(total_population = sum(poptotal)) |>
  arrange(desc(total_population))
state_population
```

```
# A tibble: 5 x 2
  state total_population
  <chr>      <int>
1 IL          11430602
2 OH          10847115
3 MI          9295297
4 IN          5544159
5 WI          4891769
```

```
state_population |>
  mutate(
    population_prop = total_population / sum(total_population)
  ) |>
  arrange(desc(population_prop))
```

```
# A tibble: 5 x 3
  state total_population population_prop
  <chr>      <int>           <dbl>
1 IL          11430602        0.272
2 OH          10847115        0.258
3 MI          9295297         0.221
4 IN          5544159         0.132
5 WI          4891769         0.116
```

Illinois is the most populated with 27% of the Midwestern population living there.

Question 8

```
state_poverty <- midwest |>
  group_by(state) |>
  summarise(mean_percbelowpoverty = mean(percbelowpoverty))
state_poverty
```

```
# A tibble: 5 x 2
  state mean_percbelowpoverty
  <chr>          <dbl>
1 IL            13.1
2 IN            10.3
3 MI            14.2
4 OH            13.0
5 WI            11.9
```

```
state_poverty |>  
  arrange(mean_percbelowpoverty)
```

```
# A tibble: 5 x 2  
  state mean_percbelowpoverty  
  <chr>          <dbl>  
1 IN            10.3  
2 WI            11.9  
3 OH            13.0  
4 IL            13.1  
5 MI            14.2
```

Indiana has the lowest average percentage below poverty across its counties, while Michigan has the highest average percentage below poverty.

Part 2

Question 9

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df
```

```
# A tibble: 5 x 3  
  var_1 var_2   var_3  
  <dbl> <chr>  <chr>  
1     10 Pizza   Apple  
2     20 Burger  Apple  
3     30 Pizza   Pear  
4     40 Pizza   Pear  
5     50 Burger  Banana
```

a.)

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df |>  
  arrange(var_2)
```

```
# A tibble: 5 x 3  
  var_1 var_2   var_3  
  <dbl> <chr>  <chr>  
1     20 Burger Apple  
2     50 Burger Banana  
3     10 Pizza  Apple  
4     30 Pizza  Pear  
5     40 Pizza  Pear
```

The code changed the order of the var_2 column from “Pizza, burger, pizza, pizza, burger” to “Burger, burger, pizza, pizza, pizza”. In short, the arrange sorted the rows in variable 2 by ascending alphabetical or numeric order.

b.)

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df |>  
  group_by(var_2)
```

```
# A tibble: 5 x 3  
# Groups:   var_2 [2]  
  var_1 var_2  var_3  
  <dbl> <chr> <chr>  
1     10 Pizza  Apple  
2     20 Burger Apple  
3     30 Pizza  Pear  
4     40 Pizza  Pear  
5     50 Burger Banana
```

The “group_by” didn’t seem to have an effect on variable 2 as the values seemed to reflect those found in the original tibble table. The “group_by” grouped the data set based on the values within the variable 2 row, but left all the rows and other variables unchanged. When comparing “arrange” and “group_by”, we can see that “group_by” only effects the data groupings while leaving everything else the same compared to the “arrange” function which simply rearranges the rows.

c.)

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df |>  
  group_by(var_2) |>  
  summarize(mean_var_1 = mean(var_1))
```

```
# A tibble: 2 x 2  
  var_2   mean_var_1  
  <chr>     <dbl>  
1 Burger      35  
2 Pizza      26.7
```

The new code groups the data based on the values within the variable 2 row with “group_by”, then forms 2 rows based on the mean values of variable 1 found within each new grouping using “summarize”.

d.)

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df |>  
  group_by(var_2, var_3) |>  
  summarise(mean_var_1 = mean(var_1))
```

`summarise()` has grouped output by 'var_2'. You can override using the `groups` argument.

```
# A tibble: 4 x 3  
# Groups:   var_2 [2]  
  var_2  var_3  mean_var_1  
  <chr>  <chr>     <dbl>  
1 Burger  Apple      20  
2 Burger  Banana     50  
3 Pizza   Apple      10  
4 Pizza   Pear       35
```

The code created groupings based on variables 2 and 3 using “group_by”, then found the mean value of variable 1 found within each new grouping using “summarize”. The message states that the summarize is based on the 2nd variable and to change that you can use the “groups” argument.

e.)

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
df |>  
  group_by(var_2, var_3) |>  
  summarize(mean_var_1 = mean(var_1), .groups = "drop")
```

```
# A tibble: 4 x 3  
  var_2   var_3   mean_var_1  
  <chr>   <chr>       <dbl>  
1 Burger  Apple      20  
2 Burger  Banana     50  
3 Pizza   Apple      10  
4 Pizza   Pear       35
```

Everything in the code and output is the same as part d apart from the “.groups” which ungroups the data after the “summarize” has been completed.

f.)

```
df <- tibble(
  var_1 = c(10, 20, 30, 40, 50),
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")
)
df |>
  group_by(var_2, var_3) |>
  summarize(mean_var_1 = mean(var_1), .groups = "drop")
```

```
# A tibble: 4 x 3
  var_2  var_3  mean_var_1
  <chr>  <chr>     <dbl>
1 Burger  Apple      20
2 Burger  Banana     50
3 Pizza   Apple      10
4 Pizza   Pear       35
```

```
df <- tibble(
  var_1 = c(10, 20, 30, 40, 50),
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")
)
df |>
  group_by(var_2, var_3) |>
  mutate(mean_var_1 = mean(var_1))
```

```
# A tibble: 5 x 4
# Groups:  var_2, var_3 [4]
  var_1 var_2  var_3  mean_var_1
  <dbl> <chr>  <chr>     <dbl>
1     10 Pizza  Apple      10
2     20 Burger Apple      20
3     30 Pizza  Pear       35
4     40 Pizza  Pear       35
5     50 Burger Banana     50
```

The first pipeline made groupings based on variables 2 and 3, then found the mean values of variable 1 within each new grouping using the “summarize” argument. The “.groups” made it so that there were no groupings after the “summarize” was complete. The second pipeline still

grouped by variables 2 and 3 but instead made a new column that shows the mean values of variable 1 within each new grouping using the “mutate” argument. All groupings remained and a new column was created.